

2)

```
import pandas as pd
```

```
# Load CSV file into a DataFrame (this is a comment)
```

```
df = pd.read_csv('path/to/file.csv')
```

```
# Print the first few rows of the DataFrame (this is a comment)
```

```
print(df.head())
```

In this script, we first import the pandas library with the **import pandas as pd** statement.

Then we use the **pd.read_csv()** function to load the CSV file located at 'path/to/file.csv' and read it into a DataFrame. Finally, we use the **df.head()** function to print the first few rows of the DataFrame.

Note: Make sure to replace 'path/to/file.csv' with the actual file path to your CSV file.

(BASCALLY IN THE ABOVE CODE WHERE IT IS MENTIONED AS PATH TO CSV FILE,
DOWNLOAD THE DATASET INTO YOUR LOCAL FILE AND MENTION THE DATA SET FILE PATH
IN THE CODE)

3) **import** pandas **as** pd

import numpy **as** np

import matplotlib.pyplot **as** plt

import seaborn **as** sns

from scipy

import stats **from** sklearn

import model_selection, feature_selection, linear_model **from** sklearn.linear_model **import**
LinearRegression **from** sklearn.preprocessing

import StandardScaler,MinMaxScaler,OneHotEncoder,OrdinalEncoder **from** matplotlib

import pyplot **from** sklearn.linear_model

import Ridge , Lasso,RidgeCV **from** sklearn.metrics

import mean_squared_error, r2_score **from** scipy.stats

import shapiro**from** statsmodels.stats.diagnostic

import normal_ad **from** statsmodels.stats.outliers_influence

import variance_inflation_factor

import statsmodels.api **as** sm **from** statsmodels.stats.diagnostic

import het_goldfeldquandt

import warningswarnings.filterwarnings('ignore') **from** sklearn.compose

import ColumnTransformer , make_column_transformer,make_column_selector **from**
sklearn.pipeline

import Pipeline **from** sklearn.feature_selection

import mutual_info_regression, SelectKBest,f_regression,RFE

df["normalized-losses"].fillna(df["normalized-losses"].astype("float").mean(), inplace=**True**)

df["bore"].fillna(df["bore"].astype("float").mean(), inplace=**True**)

df["stroke"].fillna(df["stroke"].astype("float").mean(), inplace = **True**)

df["peak-rpm"].fillna(df["peak-rpm"].astype("float").mean(), inplace = **True**)

df["horsepower"].fillna(df["horsepower"].astype("float").mean(), inplace=**True**)

```

category = ["make", "fuel-type", "aspiration", "num-of-doors", "body-style",
            "drive-wheels", "engine-location", "engine-type", "num-of-cylinders", "fuel-system"]fig,
axes = plt.subplots(nrows=5, ncols=2, figsize=(14, 18))axe = axes.ravel()for i, category in
enumerate(df_obj[category]):
    df_obj[category].value_counts().plot(kind="bar",
    ax=axe[i]).set_title(category)fig.show()fig.tight_layout()

```

```

sns.regplot(x = 'engine-size', y = 'price', data = df)plt.ylim(0,) # y axis starts from
zeroprint("correlation between engine-size and price:")df[["engine-size", "price"]].corr()
sns.regplot(x = 'highway-mpg', y = 'price', data = df)print('correlation between highway-mpg
and price:')df[["highway-mpg", 'price']].corr()
sns.regplot(x = 'peak-rpm', y = 'price', data = df)print('correlation between peak-rpm and
price:')df[["peak-rpm", 'price']].corr()
print('correlation between stroke and price :')
sns.regplot(x = 'stroke', y = 'price', data = df)
.df[["stroke", 'price']].corr()
stdscaler=StandardScaler()numeric_features = ["symboling", "normalized-losses", "wheel-
base", "length", "width",
        "height", "curb-weight", "engine-size", "bore", "stroke", "compression-
ratio", "horsepower",
        "peak-rpm", "city-mpg", "highway-
mpg"]num_transformed=stdscaler.fit_transform(X_numeric)num_transformed=pd.DataFram
e(num_transformed, columns=numeric_features)num_transformed.head()
def residual(model, feature, label):
    pred=model.predict(feature)
    df2=pd.DataFrame({ "Actual":label, "Predicted":pred})
    df2["Residuals"]=df2["Actual"]-df2["Predicted"]
    return df2
def linear_assumption(model, features, label):
    print('Linearity Check:', '\n')

    print('Checking with a scatter plot of actual vs. predicted.',
        'Predictions should follow the diagonal line.')
    df_results = residual(model, features, label)

```

```

sns.lmplot(x='Actual', y='Predicted', data=df_results, fit_reg=False, height=7)
# sns.residplot(x="Actual", y="Predicted", data=df_results)
# line_coords = np.arange(df_results.min().min(), df_results.max().max())
line_coords = np.arange(df_results.iloc[:,0:2].min().min(),
df_results.iloc[:,0:2].max().max())
plt.plot(line_coords, line_coords, # X and y points
        color='darkorange', linestyle='--')
plt.title('Actual vs. Predicted')
plt.show()

```